

Movie projections using Relational Databases and Classification

The Wonderbolts

Chris Arcand, Rich Jeffery, and Brian Salter

Higher. Faster. Further.

Motivation

Using existing data, can we accurately predict how well upcoming movies will do?

(And will this make us lots of \$\$\$?)

Motivation

Using existing data, can we accurately predict how well upcoming movies will do?
(And will this make us lots of \$\$\$?)

**Known metrics, such as:
Cast, Director, Release date, Rating,
Duration, Subject matter...**

Motivation

Using existing data, can we accurately predict how well upcoming movies will do?

(And will this make us lots of \$\$\$?)

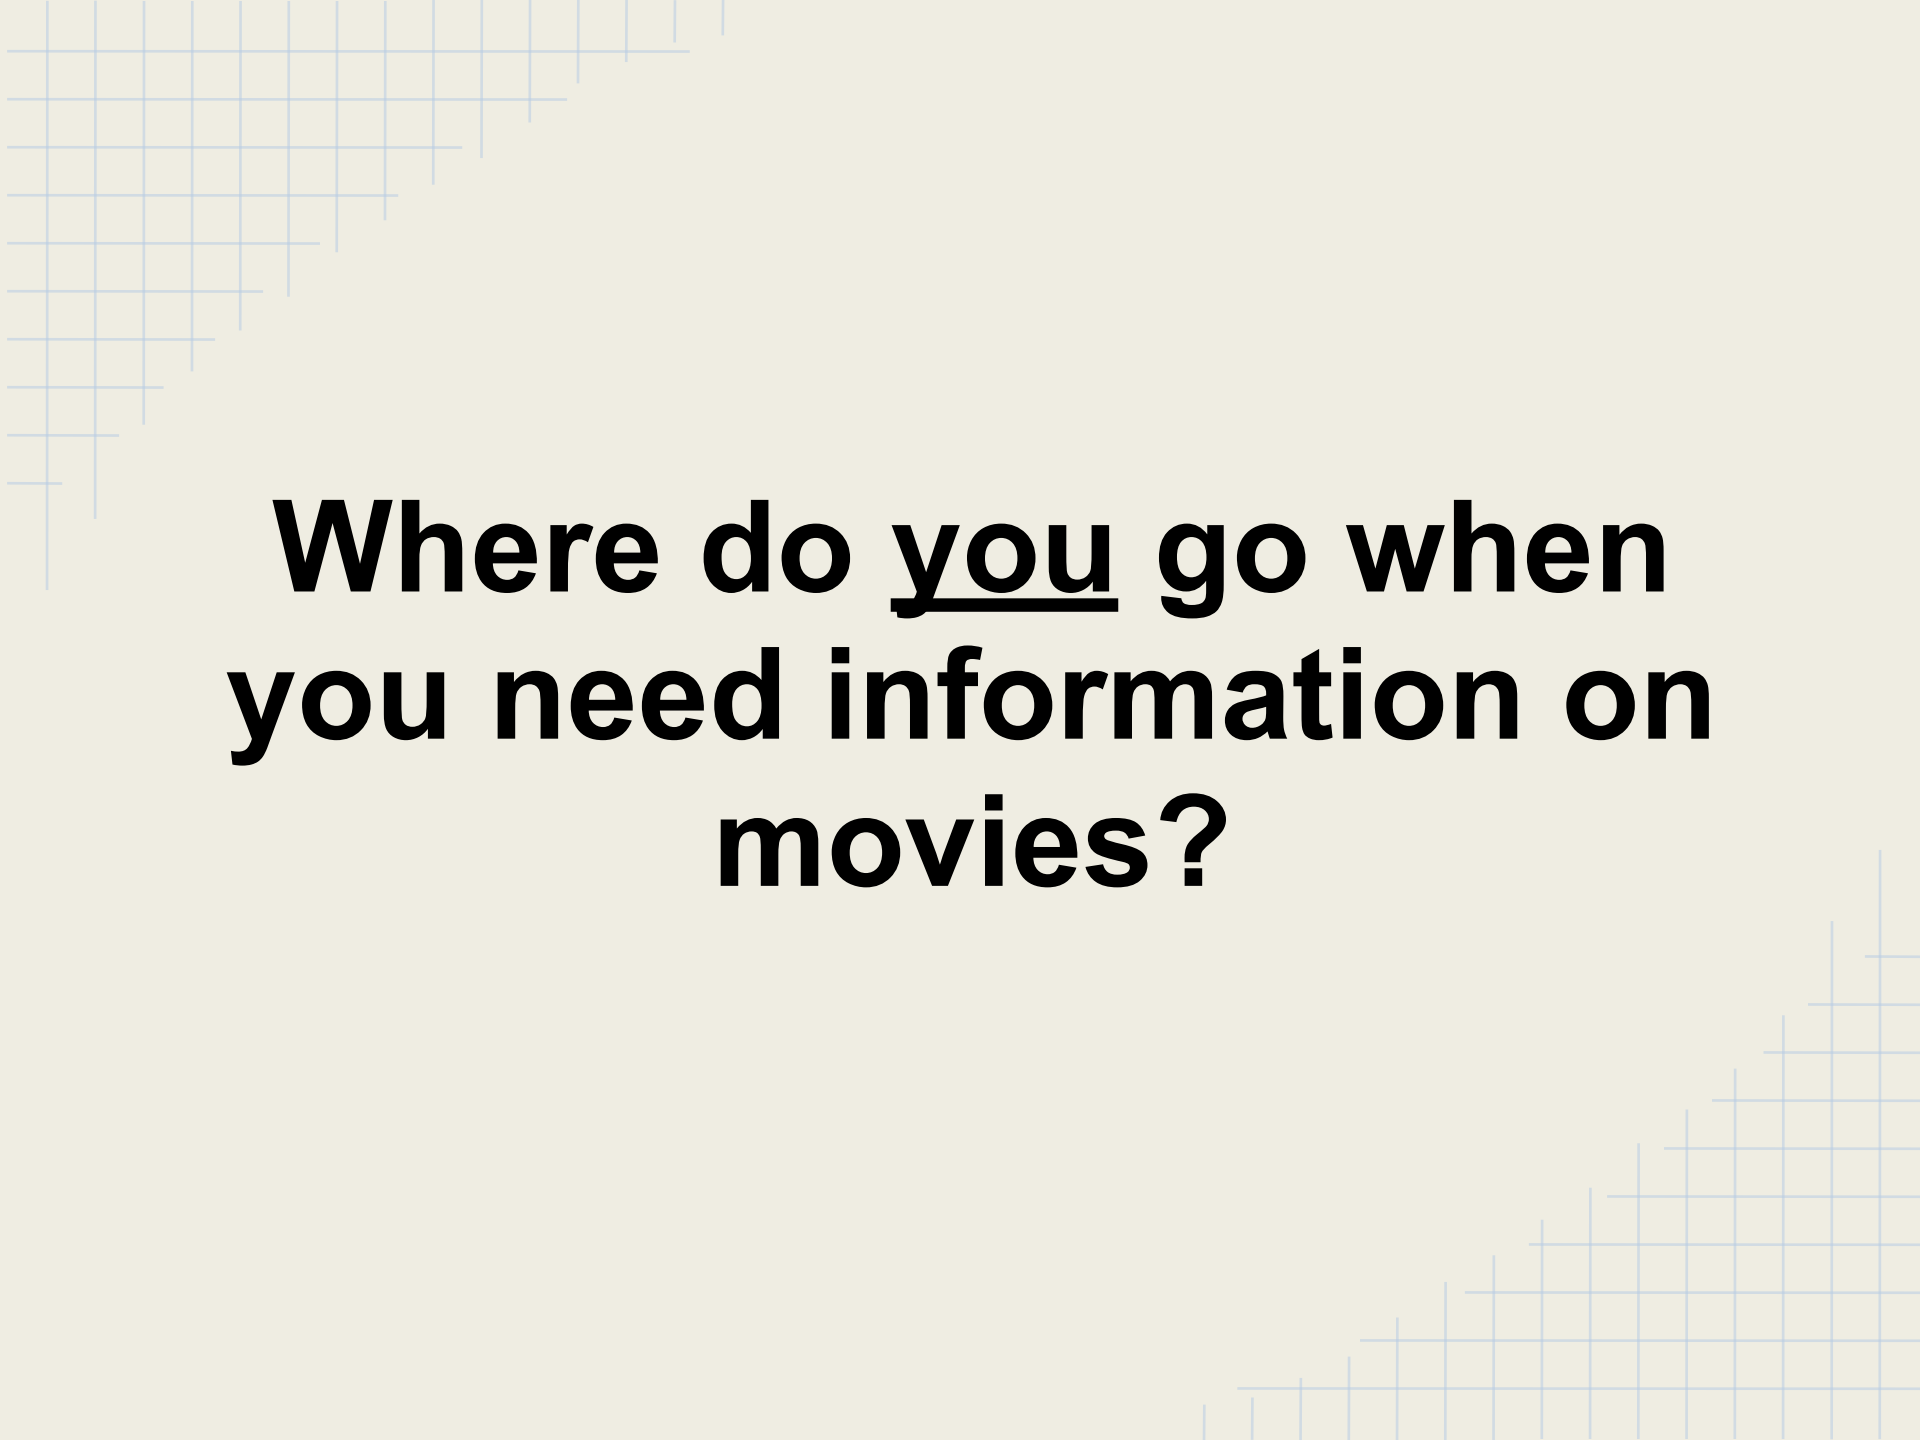
Known metrics, such as:

Cast, Director, Release date, Rating, Duration, Subject matter...


Find a pattern in these metrics, and use it to make a prediction

CHALLENGE ACCEPTED!!!





**Where do you go when
you need information on
movies?**



**Follow up question:
When you use Google,
what is the first link that
appears?**

IMDb

Internet Movie Database

The Avengers (2012) - IMDb

www.imdb.com/title/tt0848228/?ref=sr_1

Google

Unlimited access to thousands of movies
Prime Instant Video


JUSTIFIED thegoodwife THE SHIELD

amazon Sign Up

IMDb Find Movies, TV shows, Celebrities and more... All

Register | Login | Help

Movies TV News Showtimes Community IMDbPro Apps Your Watchlist

The Avengers (2012) 

PG-13 143 min - Action - 4 May 2012 (USA)

8.3 Your rating: ★★★★★★ -/10
Ratings: 8.3/10 from 494,735 users Metascore: 69/100
Reviews: 1,429 user | 610 critic | 43 from Metacritic.com

Nick Fury of S.H.I.E.L.D. brings together a team of super humans to form The Avengers to help save the Earth from Loki and his army.

Director: Joss Whedon
Writers: Joss Whedon (screenplay), Zak Penn (story), 1 more credit »
Stars: Robert Downey Jr., Chris Evans, Scarlett Johansson | See full cast and crew

+ Watchlist Watch Trailer Share...

amazon Watch now At Amazon Instant Video

amazon Own it Buy it at Amazon.com

Top 250 #149 | Nominated for 1 Oscar. Another 14 wins & 56 nominations. See more awards »

Watch these movies and thousands more on Prime Instant Video

amazon Sign Up

ad feedback

Quick Links

Full Cast and Crew Plot Summary
Trivia Parents Guide
Quotes User Reviews
Awards Release Dates
Message Board Company Credits

Explore More

Like 114,943 people like this. Be the first

How to get the data?

IMDB API

A lightweight, RESTful web interface to IMDb
www.imdbapi.com

(Make JSON queries to fill our database)

How to get the data?

~~IMDB API~~

~~A lightweight, RESTful web interface to IMDb~~
~~www.imdbapi.com~~

- Lots of overhead in automating queries
- Huge data. Bandwidth isn't cheap (or plentiful)
 - 3,000 queries per hour limit
- 4,115,950 rows in 'person' table -> 57 days of processing for ONE table



IMDB API would have been slow. Very, very slow.

How to get the data?

IMDbPY

Open source project to retrieve and manage IMDb data
imdbpy.sourceforge.net

- Written in pure Python (with a few lines of C!)
- Released under GPL 2 License
- Can both retrieve data from IMDb directly (web)
- Can also parse data in IMDb's plain text files

How to get the data?

IMDbPY

Open source project to retrieve and manage IMDb data
imdbpy.sourceforge.net

- Written in pure Python (with a few lines of C!)
- Released under GPL 2 License
- Can both retrieve data from IMDb directly (web)
- Can also parse data in IMDb's plain text files**

Minor forgotten detail...

IMDb releases subsets of the plain text file
dumps of their database!
www.imdb.com/interfaces/

49 files
6.75GB when uncompressed

Part 3/3 - The Database

-IMDbPY doesn't support Oracle
(open source, yo)

Suggestion: SQLite -> Oracle

Part 3/3 - The Database

-IMDbPY doesn't support Oracle
(open source, yo)

~~**Suggestion: SQLite -> Oracle**~~

-SQLite doesn't work well with large data

-In addition, 4GB limit on Oracle tablesize == bad

MySQL to the rescue

Finally, loading the data

IMDb text files + IMDbPY + MySQL


~1.5 hours to process in final run

Many attempts required to get it right:

- Database issues**
- Dependency issues within Python**
- Resource issues (VPS vs. decent machine)**


Lesson learned: Big data requires thoughtful planning and knowledge to set up and maintain effectively.

What does the data look like?



Let's see an example!






What if we wanted to actually do something with this data as is? Suppose we wanted to find out which actor played the best James Bond.



First things first...

Firstly we need to find the `role_id` of James Bond from the `char_name` table.
Easy enough.

```
1 • SELECT * FROM imdb.char_name WHERE name like '%James Bond%';
```

Filter:		↔	Edit: 		File: 		Autosize: 
	id	name	imdb_index	imdb_id	name_pcode_nf		
	14376	Himself - Production Designer for the James Bond Films, 1962-1979	NULL	NULL	H5241		
▶	48229	James Bond	NULL	NULL	J5215		

Now what?

Well now we need to find the `person_id` for every person who has played James Bond. We do this by looking in the `cast_info` table for the `role_id`.

```
1 • Select * FROM imdb.cast_info where (person_role_id = 48229 and (role_id = 1 or role_id = 2));  
2
```

id	person_id	movie_id	person_role_id	note	nr_order	role_id
170773	26494	1663805	48229	NULL	14	1
185006	29068	1969431	48229	NULL	NULL	1
684727	96475	2266884	48229	(as Ali G)	1	1
773902	108370	1886466	48229	(voice)	NULL	1
831230	116515	1834552	48229	NULL	59	1
895478	124908	2385540	48229	(voice)	NULL	1
988634	138066	1600995	48229	NULL	1	1
1048662	146780	1600993	48229	(voice)	NULL	1
1048663	146780	1600995	48229	(voice)	1	1
1048675	146780	2399077	48229	(voice)	3	1
1058788	147915	50674	48229	NULL	NULL	1

...Many queries later...

- Once we have the person_id list for people that played James Bond, we need to get scores for each movie where that actor played James Bond and find their average review score.
- We also need to limit our search to theatrical releases only.
- This involves a query that joins role_id, cast_info, title, movie_info_idx, and info_type tables...
- After taking half an hour to run the query we get the following data:

Actor	Average James Bond Film Score
Sean Connery	7.0
George Lazenby	6.8
Roger Moore	6.6
Timothy Dalton	6.6
Pierce Brosnan	6.5
Daniel Craig	7.5

What if we wanted to do something simpler?

What if we just wanted to decide how good an actor is by examining the average scores of his movies?

The query for this question looks like:

```
Select AVG(info) FROM (Select imdb.cast_info.movie_id, info, info_type_id, person_id, role_id
FROM imdb.cast_info inner join imdb.movie_info_idx on imdb.cast_info.movie_id = imdb.
movie_info_idx.movie_id) As T where (person_id = 1468525 and (role_id = 1 or role_id = 2) and
info_type_id = 101);
```

Running this query takes upwards of 10 minutes and tells us that the late Rod Steiger has an average review score of 6.2.



That seems like a lot of time and work...

If we want to look at hundreds of thousands of movies, and millions of actors we need to change something...



We need to change the world... or at least the structure of our database.

Data Manipulation

- 2.7 million columns
- 100K rows

How do we deal with this much data?



What we lost...

- Actors not in a movie
- Redundant entries
- Independent films
- Low-budget films
 - < \$1,000,000




dataGet.py

- Matlab hates us
- Store data in text files
 - data_2000.txt
 - target_2000.txt
- Read in data

Perceptron

- Many inputs
- Many weights



How it works...

- Dot products
 - Feature vector
 - Weight vector

$$f(x) = \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

- Minimize test error
 - misclassified / total

Naive Bayes Classifier

- Probabilities
 - $P(\text{Class}|\text{Features})$
- Learns $P(F|C)$

$$p(C|F_1, \dots, F_n) = \frac{1}{Z} p(C) \prod_{i=1}^n p(F_i|C)$$

Computes $P(C|F)$ to predict

6.5 Star Threshold

```
Brians-MacBook-Pro-2:Desktop salterbw$ python doTests.py  
2011  
2012  
Testing error: 0.349799732977  
Naive Bayes Error: 0.427236315087  
Brians-MacBook-Pro-2:Desktop salterbw$
```


7 Star Threshold

```
Brians-MacBook-Pro-2:Desktop salterbw$ python doTests.py
2011
2012
Testing error: 0.217623497997
Naive Bayes Good Error: 0.229639519359
Naive Bayes Bad Error: 0.229639519359
Brians-MacBook-Pro-2:Desktop salterbw$
```

What comes next...

1. Increase Timespan
2. More Features
3. Multiclass
4. Multi-layers

Summary

1. Big data == Lots of time to load

Summary

1. Big data == Lots of time to load
2. Big data == Lots of time to maintain

Summary

1. Big data == Lots of time to load
2. Big data == Lots of time to maintain
3. Big data == Lots of time to analyze
4. Make good, smart choices about how/what you are doing
5. Challenge defeated, @ArtsJournal

FIN

